Responsible AI, Agentic AI, and Prompt Skills Development for IA

Event

Event Date

The better the question. The better the answer. The better the world works.

Shape the future with confidence

Learning Objectives for this course

1. Apply Responsible AI Principles in Auditing:
   - Participants will explore frameworks guiding the responsible development and deployment of AI, focusing on fairness, bias mitigation, and accountability. They will learn how to integrate these principles into auditing practices to ensure ethical and compliant use of AI technologies.

2. Evaluate Agentic AI in Business Contexts:
   - Attendees will analyze real-world applications of Agentic AI across industries, assess associated risks—including operational, strategic, and compliance—and understand its influence on organizational decision-making, transparency, and accountability.

3. Develop Effective Prompting Skills for Generative AI in Auditing:
   - Participants will gain hands-on experience designing prompts tailored to audit scenarios, enhancing their ability to elicit accurate and relevant responses from Generative AI tools throughout the audit lifecycle.

Agenda

Module Name	Timing
Responsible AI Principles in Auditing	8:30 - 9:20 AM
Recognize Agentic AI Applications in Business	9:20 - 10:10 AM
Break	10:10 - 10:30 AM
Demonstrate Prompting Skills in GenAI Tools for Audit Applications	10:30 AM - 12:00 PM

Module 1: Responsible AI Principles in Auditing

Event

The better the question. The better the answer. The better the world works.

Shape the future with confidence

Key Responsible AI Questions

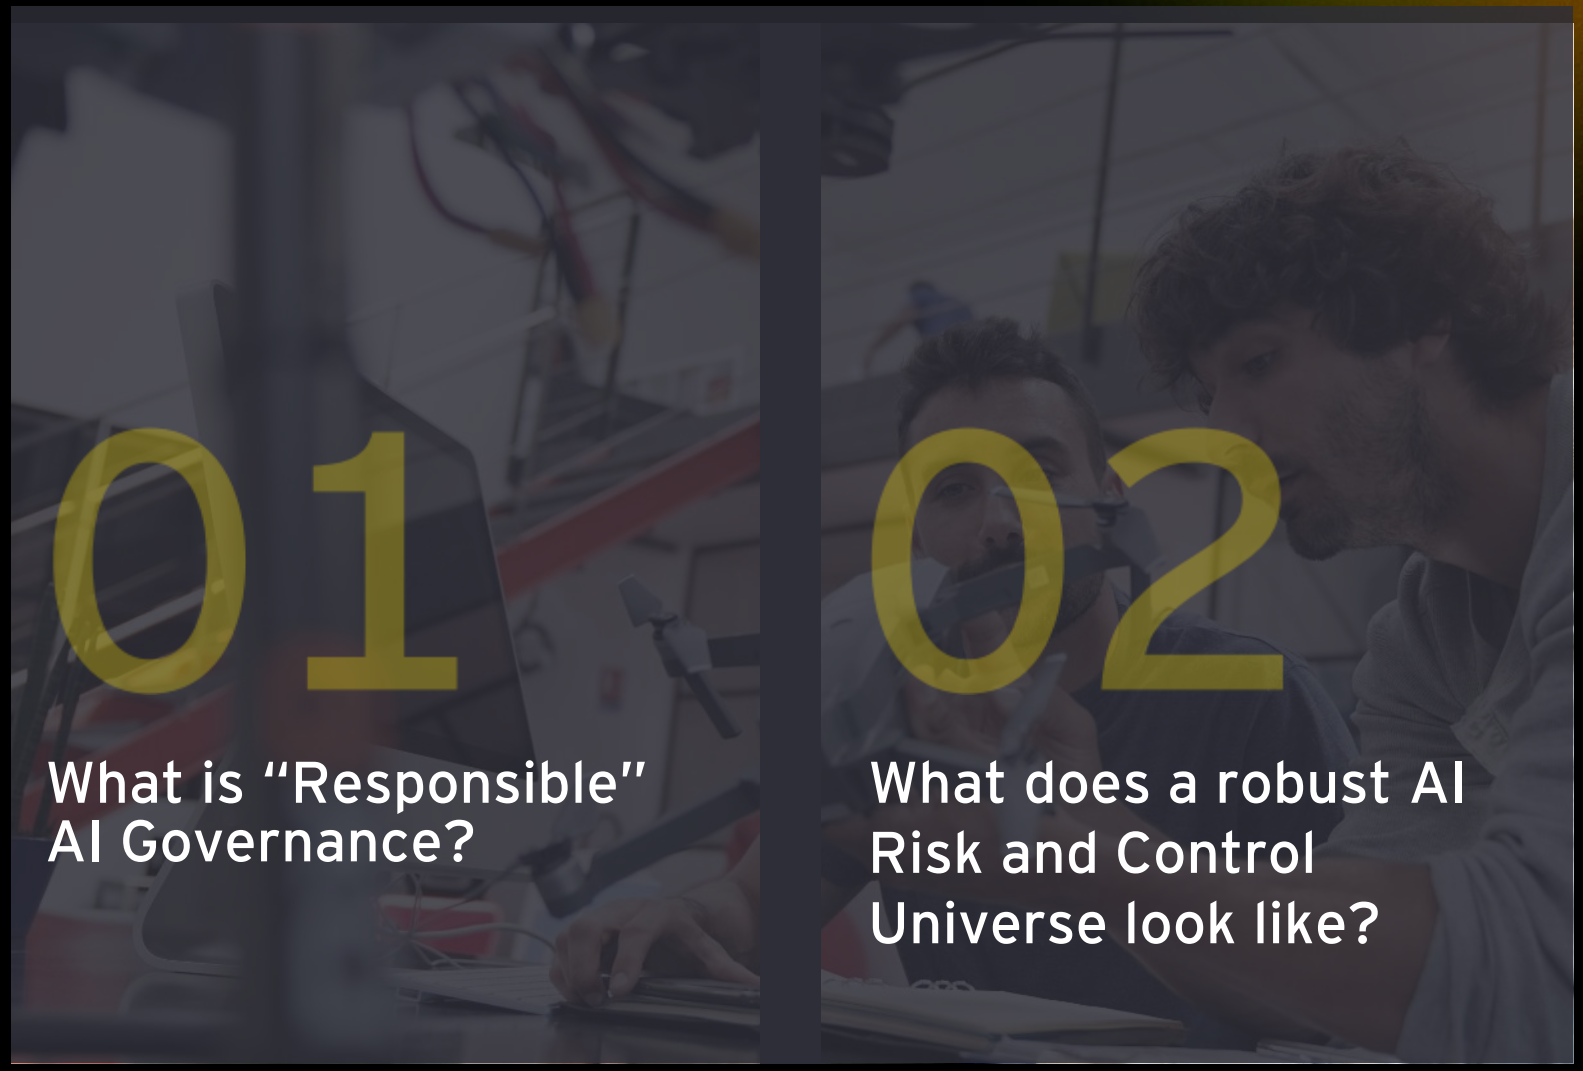

01

What is "Responsible" AI Governance?

02

What does a robust AI Risk and Control Universe look like?

Discussion Point

In 1-2 words, what benefits do you see in having a strong Responsible AI/AI Governance program?

Please navigate to the Microsoft Whiteboard to enter your response!

We are in an AI revolution that will introduce unprecedented change to how we define opportunities and risks

The scale, velocity, and uncertainty of change, combined with the ubiquitous and ever-evolving risk landscape of widespread AI adoption, raises a call to action around risk governance and management. This call to action is enterprise-wide, transformative to legacy lines of defense, and demands new processes, technology enablement, and capabilities beyond what narrow AI adoption required. It will also require a continuous recalibration of your Responsible AI (RAI) risk management strategy to align with evolving RAI priorities around enabling competitiveness, compliance, and protection.

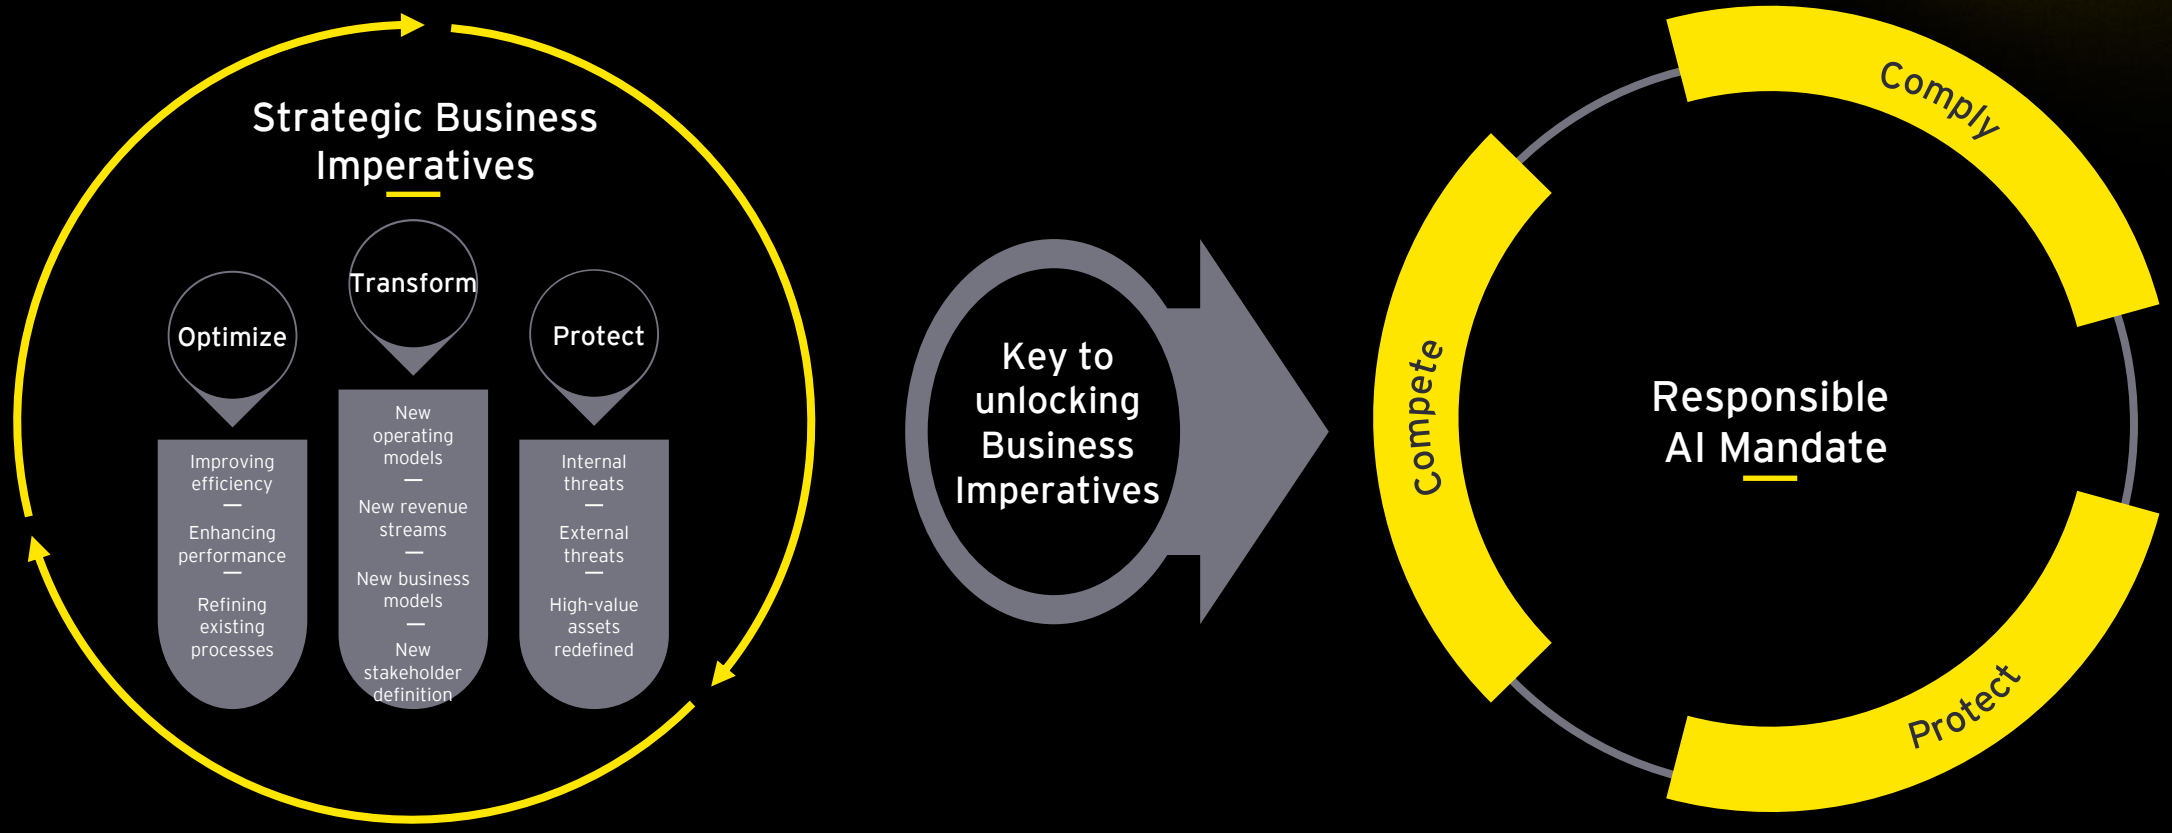

Responsible AI adoption should start with governance

By setting up the right governance framework for AI, an organization can amplify the positive outcomes of innovation and mitigate potential risks, especially considering the rapid pace of technological innovation in this field.

Organizations who

Lead with governance

- Understand and mitigate the risks of AI before they materialize
- Proactively identify high-value areas for AI deployment
- Develop direct value generation through clear and consistent innovation
- Scale learnings and leading practices discovered in early pilots
- Accelerate ROI by protecting revenue in addition to generating it
- Adopt AI at scale and maximize impact across functions and BUs
- Have a focused and consistent approach to innovation and scaling

Organizations who

Do not lead with governance

- Adopt AI in silos which limits scale, ROI, and visibility
- Create prioritization chaos, which results in fragmented investment
- Reactively address risks, damaging both reputation and finances
- Repeat mistakes across BUs throughout the AI adoption lifecycle
- Result in slow-paced innovation, culture regression and reduced collaboration

EY

Essential elements of your Responsible AI framework

Why does your AI need to be responsible? And what does "responsible" mean?

1 Responsible AI strategy that drives business performance

Purpose — the **why?**

Compete – drive competitive advantage with AI
Protect – safeguard the company's assets and preserve value with AI
Comply – achieve speed to market while complying with applicable regulations

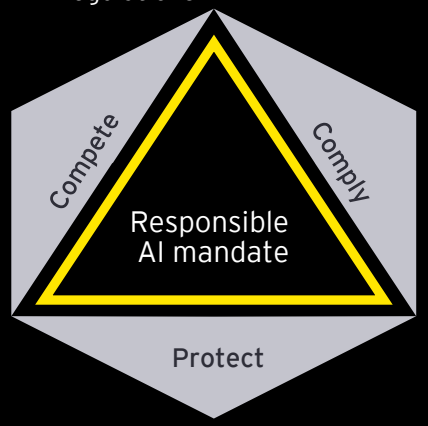

2 Aligning your RAI strategy with your values

Value — **the what?**

3 Optimize your AI assets using the RAI Framework

Principles — the **how?**

Delivered through the enablement of the following principles and measure performance through their **confidence index**

- Accountability
- Explainability
- Reliability
- Fairness
- Transparency
- Security
- Compliance
- Sustainability
- Data protection

4 Driving success

Navigate — **the governance**

5 Maintain consistent performance

Measure — Performance management

EY

Responsible AI principles

Compliance: Adherence to relevant laws, regulations, professional standards and economics.

Sustainability: Consideration of the impacts technology has on promoting physical, social, economic and environmental wellbeing.

Data protection: Protection of data in AI systems to maintain consistency and confidentiality of business and employee personal information.

Accountability: Unambiguous ownership over AI systems.

Explainability: Appropriate levels of explanation are enabled to reasonably understand, challenge and validate the decision criteria and output of AI systems.

Reliability: AI systems can continually preform at desired level with reproducibility, precision and accuracy.

Fairness: AI models are free from bias and their methodology can be explained.

Transparency: Appropriate levels of disclosure over the purpose, design and impact of AI systems.

Security: Establishes that the AI system works as intended and is protected from unauthorized access.

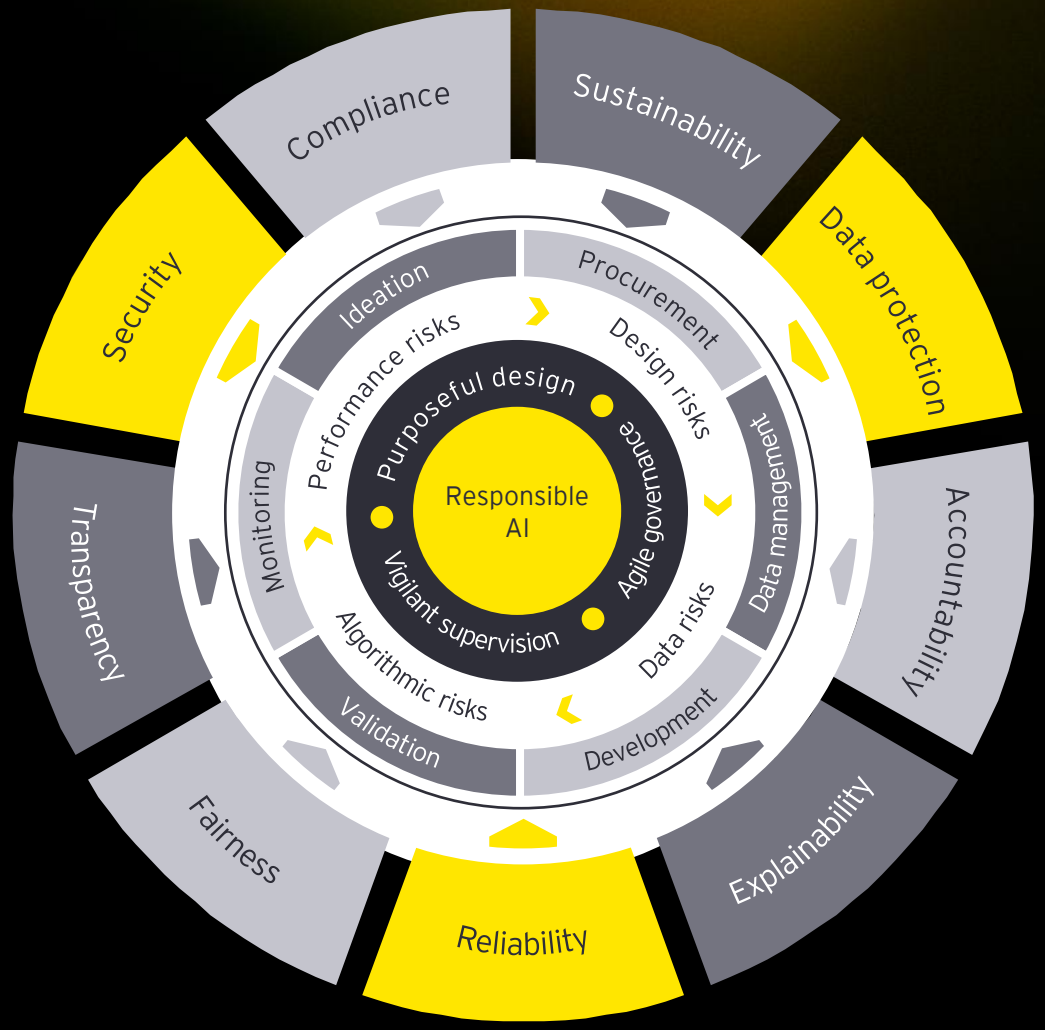

Responsible AI (RAI) Lifecycle - Key Considerations

01 | RAI fundamentals

A
AI system definitions

B
AI inventory

C
Responsible AI impact assessment

02 | RAI model development lifecycle

A. AI model development lifecycle iterative sprint(s)

- A.1 Business case assessment
- A.2 Data acquisition
- A.3 Data prepareration
- A.4 Model selection or feature engineering
- A.5 Model training
- A.6 Model testing and evaluatoin
- A.7 Model validation and validation
- A.8 Model deployment
- A.9 Ongoing monitoring

"TEVV"

- B. Fairness and equity
- C. Third-party risk assessment
- D. Resiliency or business continuity
- E.1 Security measures
- E.2 Incident management
- E.3 Data leakage and exposure
- E.4 Logging and monitoring

03 | RAI transversal requirements

- A. Policies and procedures
- B. Training and communication
- C. Transparency
- D. Risk assessment
- E. Sustainability and environment
- F. Notifications to authority

EY

Discussion Point

You are auditing an AI tool in-use. After asking questions on the tool's output generation logic, you realize that the tool owners have challenges with describing the steps the tool takes in generating output, and they struggle to showcase how tool outputs are validated. What Responsible AI pillar is not being demonstrated?

Please navigate to the Microsoft Whiteboard to enter your response!

EY

Discussion Point

You are auditing an AI tool in-use. You observe that prompts input into the tool often return information that is identifiable to a company employee, including their email and home address. The tool is accessible by all employees within the organization. What Responsible AI pillar is not being demonstrated?

Please navigate to the Microsoft Whiteboard to enter your response!

EY

AI Risk Origination Points

Origination Points

1. Use of third-party AI solutions including public AI tools
2. AI embedded into tools and services from routine vendors and service providers
3. Acquisitions of companies developing or using AI

Example actions to manage risks

1. Implement written standards, policies, procedures and protocols to acquire, use and rely on third-party data and third-party AI systems (1, 2, 3)
2. Create an intake process and risk-ranking criteria for AI solution providers to allow for inherent and residual risk scoring (1)
3. Update existing vendor questionnaires and vendor risk management frameworks to identify AI integration by existing service providers (1, 2)
4. Integrate ethical AI and TPRM programs into M&A due diligence process. (3)
5. Create third-party AI inventories, incorporating residual risk-ranking and associated monitoring, testing or auditing requirements (1, 2)
6. Enhance your contracts by adding specific terms, like: (1, 2)
   a. "right to audit clauses" – to be able to inspect and test vendors' models
   b. "collaboration clauses" – to ensure cooperation from vendors in the context of regulatory inquiries and investigations
   c. "Indemnity clauses" – to limit liability exposure from vendors

EY

RAI Risk Universe

Responsible AI

How are you defining your AI risk environment?

Considering risk from different pillars and lens is necessary for holistic risk management of the AI Program: Strategic lens, Operations lens, Model lens, Regulatory lens.

Common AI Applications in Business and Operations

	AI operational risks	AI governance risks	AI model risks	AI strategic risks
AI-enabled services Traditional services enhanced with AI capabilities to improve efficiency, accuracy and customer experience	▸ Integration complexities ▸ Data security ▸ Legacy system incompatibility ▸ Change management	▸ Inconsistent AI control ▸ Implementation ▸ Absence of clear liability ▸ Centralized AI risk visibility	▸ Data bias ▸ Inaccurate predictions ▸ Feedback loop vulnerabilities ▸ Adversarial attacks	▸ Misaligned AI objectives to business goals ▸ Ethical oversights ▸ Market perception ▸ Partnership dependencies
AI as a service (AIaaS) Process and synthesize information like human beings and with exceptional speed	▸ Data security ▸ Misconfigured parameters ▸ Talent gaps ▸ Unintended data outputs	▸ Proprietary obfuscation ▸ Inconsistent standards ▸ AI product/use case portfolio management	▸ Model drift ▸ Spoofing ▸ Data leakage protection ▸ Training-testing mismatch	▸ Strategic implementation overreliance ▸ Scalability barriers ▸ Regulatory blind spots ▸ Intellectual property protection
Automation and monitoring Development of AI technologies to automate routine tasks and monitor systems or processes	▸ Supply chain real time processing delays ▸ Data quality ▸ Generative AI access management ▸ AI components asset management	▸ Centralized AI risk visibility ▸ Inconsistent AI control implementation ▸ Absence of clear liability	▸ Overfitting/underfitting ▸ Hidden layer ambiguities ▸ Model poisoning ▸ Feature redundancy	▸ Misallocation of resources ▸ Ethical oversights ▸ Geopolitical challenges ▸ Late adoption competitive disadvantage
Operational efficiencies Optimizing business operations through automation, resource optimization and supply chain management	▸ Misguided asset allocation leading to liquidity issues ▸ Change management ▸ Data quality ▸ Reproducibility issues	▸ Inconsistent standards ▸ AI product/use case portfolio management ▸ Proprietary obfuscation	▸ Data bias ▸ Inaccurate predictions ▸ Data fluctuation distortion ▸ Information integrity	▸ Misaligned AI objectives to business goals ▸ Scalability barriers ▸ Strategic stagnation ▸ AI technical talent attrition
Compliance monitoring Ensuring that business operations adhere to legal and regulatory standards and automating reporting processes	▸ Unintended data outputs ▸ Legacy system incompatibility ▸ Data security ▸ Inadequate infrastructure	▸ Inconsistent AI control implementation ▸ Absence of clear liability ▸ Centralized AI risk visibility	▸ Adversarial attacks ▸ Data leakage protection ▸ Spoofing ▸ Feedback loop vulnerabilities	▸ Ethical oversights ▸ Regulatory blind spots ▸ Strategic stagnation ▸ Market perception

EY

RAI Governance Current State Discovery Approach – Illustrative

Defining the scope

AI Inventory
List of categorized AI-related assets

AI Policy & Design standards
Governing principles and design criteria for AI solutions

Resource Management
Allocation and optimization of resources for AI projects

Risk & Control Framework
Methodology for identifying, assessing, and mitigating risks associated with AI

Data Management
Strategies for handling data acquisition, storage, processing, and security

Secure Architecture
Focus on cybersecurity and data protection throughout the design of AI solutions

Compliance & Legal
Adherence to legal standards and regulatory requirements

Ethics and Governance
Oversight to ensure AI is used in an ethical manner, beyond policy and design standards

Operating model
Operating model definition, business model alignment, roles and responsibilities clarity

Performance Metrics
How AI success and impact are measured

AI Training & Development
Ongoing education and development programs for staff involved in AI projects

Innovation and Research
Dedicated efforts for AI technology development and procurement

Stakeholder Engagement
Involvement of internal and external stakeholders in the AI program

AI Strategy & Roadmap
Long-term planning and strategic direction of AI initiatives

Legend

- Controls are typically at organization level
- Controls are typically at system level
- Controls are typically at both system and organization level

Select assessment frameworks

- NIST
- Marketplace benchmarking
- EU AI Act
- Maturity model

EY

AI Governance Assessment

Compliance

Regulatory requirements

Authoritative sources

Business context

Control Objectives

Step 1

Enterprise Level

Control

Control

Control

Control

Step 2

Engineering

Enablers and templates to provide additional clarity on how enterprise level controls can be consistently performed by the engineering or products team.

Example:

1. Model card template
2. AI use case level risk assessment approach & templates
   a. define requirements for relevant RAI principles, TEVV
   b. Documentation requirements

Step 3

Developer's Guide to RAI

AI Use Case Level

Control Activities

Control Activities

Control Activities

Control Activities

Control Activities

Control Activities

Step 4

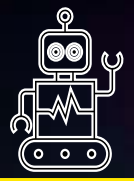

Operationalize RAI via high degree of technology integration

Step 5

Discussion Point

What role do you see IA playing when it comes to Responsible AI?

Is IA already performing some of those actions?

Please navigate to the Microsoft Whiteboard to enter your response!

EY

What capabilities and enablement do IA teams need?

Transforming internal audit for AI-driven complexity: new skills, strategic partnerships, and evolving roles

Revised audit methodologies

- Develop methodologies for auditing AI algorithms, including validation of model accuracy, fairness, and performance over time.
- Use scenario analysis to assess the impact of different risk scenarios on AI systems and incorporate these insights into audit plans.

AI risk assessment

- Early engagement with AI teams to understand objectives, design and potential risks
- Use structured AI Risk and Governance Assessment frameworks
- Regulatory Mapping
- Audit work-programs and playbooks
- Definition of risk, impact, KPIs and related performance metrices
- Facilitate joint risk assessment workshops for alignment and insights

Engage with external experts

- Collaborate with external experts, consultants, and industry groups to stay informed about the latest trends, best practices, and emerging risks in AI.
- Regularly benchmark IA practices against industry standards and leading practices to ensure continuous improvement and relevance.

Enabling tools and infrastructure

- Utilize smart audit software with automated workflows; built in work-programs and playbooks; supporting automated testing and reporting
- Enhance reporting capabilities – interactive dashboards
- Introduce model evaluation tools and reports
- Integrate AI-powered tools into IA's own processes - predictive analytics for identifying emerging risks, automated risk assessment platforms, etc.

Automated analytics and reporting

- Automated monitoring of action plans
- Customized report generation based on templates
- Deep dive into performance with visual analytics
- Trend analysis of past audits

Enhance communication and reporting

- Enhance communication with functions like General Counsel's Office (GCO) and Engineering to stay updated on regulatory changes and ensure compliance.
- Form a cross-functional committee and hold regular meetings or set up dedicated channels and reporting mechanisms.

EY

How are IA functions responding?

- Seat at the table around AI governance
- Performance of AI framework and governance audits
- Performance of AI system audits and product reviews
- Contribute to raising the Enterprise IQ around Responsible AI

The minimum questions to ask if we encounter AI in today's audits

Just as with any IT system, AI systems require initial identification and general auditing inquiries, such as determining the system's owner, assessing whether there are existing controls, and identifying the associated risks, before delving into more detailed questions specific to AI.

Questions for Management

1. Can you describe the primary functions of the AI systems in our operations?
2. How is the AI system integrated into the current business processes?
3. What data source does the AI system use, and how is data quality assured?
4. How have you evaluated the risks associated with the AI system, including ethical, privacy, and security risks?
5. How are changes in AI algorithms and models managed and communicated within the organization?
6. Are we reliant on external vendors for AI technology? How do we manage this dependence?
7. How do we ensure compliance with current and upcoming regulations specific to AI?
8. What contingency plans are in place for potential AI systems failures or breaches?
9. What ethical framework guides our use and development of AI technologies?
10. What security measure are in place, and how are errors or anomalies addressed?
11. What feedback mechanisms are in place for continuous improvement of the AI system?

Questions to think about

1. Do I have sufficient understanding of AI technology and its implication in auditing?
2. Have I identified all potential risks associated with the AI system, including ethical, privacy, and security risks?
3. Am I aware of the regulatory and compliance requirements specific practices?
4. How will the AI system impact our existing audit methodologies and practices?
5. Do I need to seek external expertise or additional training to effectively audit this AI system?
6. How can I ensure continuous monitoring and review of the AI system's performance and compliance?
7. Are there any biases inherent in the AI system that could affect audit findings?
8. What are the long-term implications of this AI system on our organization's risk profile?
9. How can I effectively communicate AI-related risks and recommendations to management and stakeholders?
10. Are the existing internal controls adequate to manage the risks presented by the AI system?

Module 2: Agentic AI Applications in Business

Event

The better the question. The better the answer. The better the world works.

Shape the future with confidence

Definition of an AI Agent

An "Agent" is software that performs tasks, makes decisions, and interacts with technology and humans, leveraging AI to achieve goals.

Agentic AI – The next wave of AI

Gartner's Top 5 Strategic Technology Trends

1. *Agentic AI*
2. AI Governance Platforms
3. Disinformation Security
4. Post-Quantum Security
5. Ambient Invisible Intelligence

Source: Technology for Trends 2025: Gartner Top 10 Strategic Technology Trends

Snowflake bares its agentic AI plans by showcasing its Intelligence platform

News
Nov 12, 2024 • 5 mins

Citizens to pair AI agents with automation to improve productivity

By Suman Bhattacharyya November 07, 2024, 3:16 p.m. EST 4 Min Read

ServiceNow partners with Nvidia to accelerate agentic AI adoption

Salesforce's new AI platform lets companies build their own digital agents to automate tasks

Matthew S. Smith Nov 11, 2024, 11:00 AM PST Share Save

Zenity Raises $38 Million to Secure Agentic AI

Agentic AI security startup Zenity has raised $38 million in a Series B funding round led by Third Point Ventures and DTCP.

Microsoft unveils open-source multi-agent AI system Magnetic-One

The system works with a lead agent called the Orchestrator which then directs the other four agents to solve tasks

Published - November 07, 2024 12:40 pm IST

IBM: Enterprise IT Facing Imminent AI Agent Revolution

Published October 16, 2024 | Written by Ben Abbott

Forecast Statistics

Day-to-Day Work Decisions
15%
made autonomously by 2028

Enterprise Software Applications
33%
will include Agentic AI by 2028

Expected CAGR
40.2%
from 2024-2032

Economic Contribution from GenAI
$2.6t+
by 2028

Sources: How Intelligent Agents in AI Can Work Alone | Gartner, Agentic AI Market Expected to Grow at 40.2% by 2032, Economic potential of generative AI | McKinsey

EY

Agentic AI – The next wave of AI

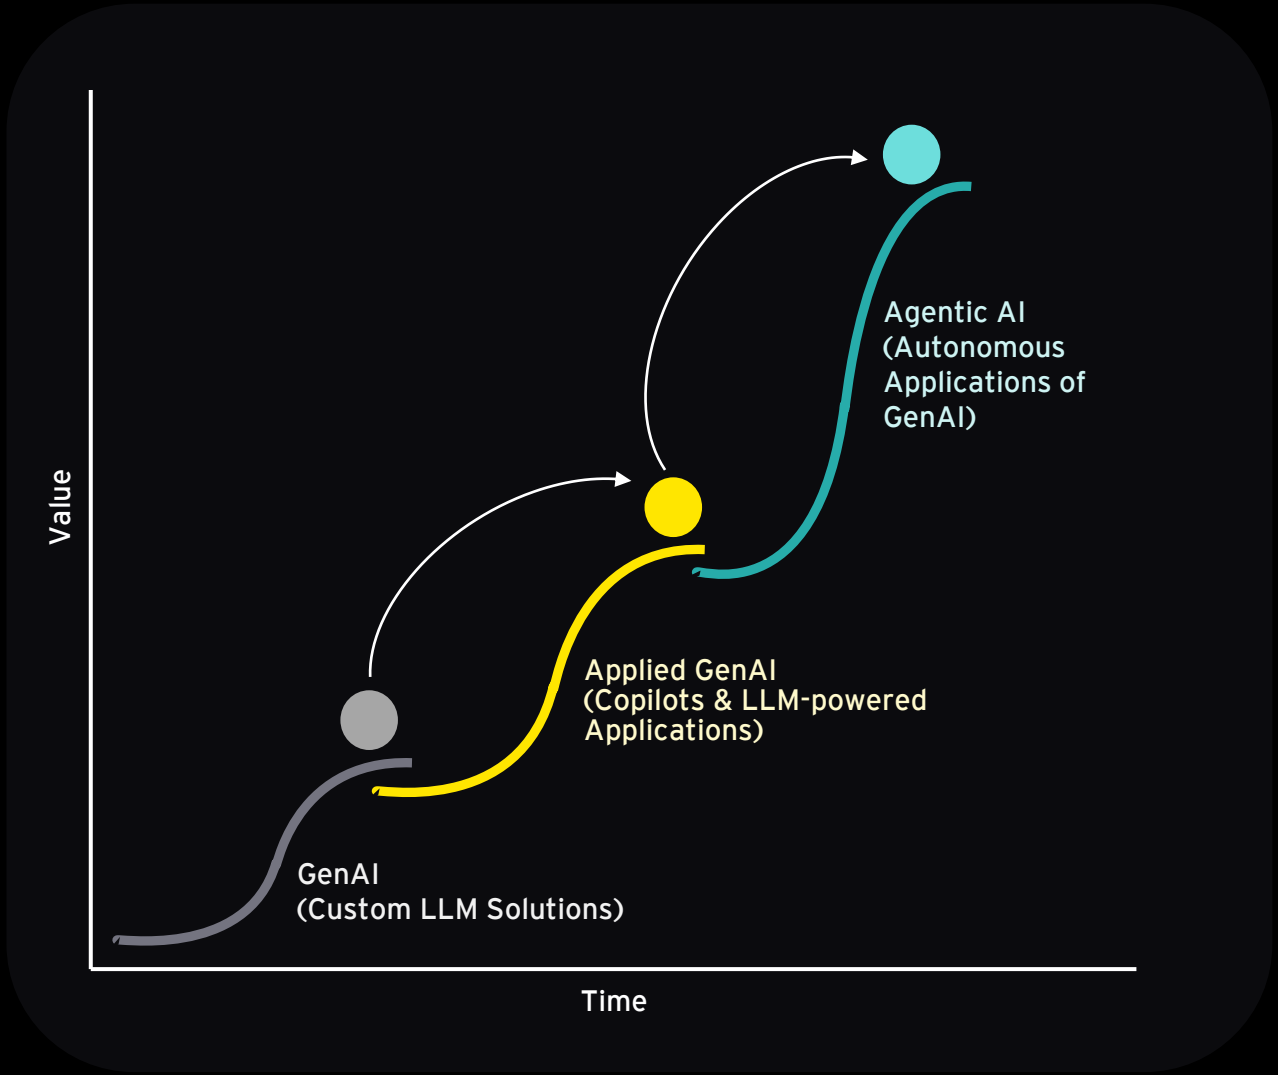

The emergence of Agentic AI represents a transformative shift towards autonomous and intelligent operations...

Enabling AI systems to independently make decisions and coordinate complex tasks...

Driving the potential for unprecedented innovation and operational reshaping across all organizations

Key Aspects of Agentic AI

Software that **autonomously perceive, act, react, and interact** with both humans and other agents alike in achieving its goals.

Perception

The ability to sense, interpret, and contextualize data from the environment, in various forms such as textual or visual inputs

Perception enables Agentic AI to gather and contextualize information to make informed decisions, allowing it to interact effectively and achieve its goals

Interaction

The ability to communicate, act with, and react to humans, computer systems, and other agents, facilitating meaningful exchanges of information and collaboration.

Interaction allows Agentic AI to work towards its objectives, enhancing its ability to achieve complex tasks and function within a broader ecosystem

Autonomy

The capability to operate independently, make decisions, and perform tasks without need for human intervention.

Autonomy allows Agentic AI to function efficiently, reducing reliance on humans and enabling it to handle complex tasks, driving innovation and productivity.

EY

Enterprise Agentic AI - Reference Architecture

Operationalizing Agentic AI at scale:

- **AI Agent Marketplace**
  Deploy domain-specific AI agents for sales, marketing, and support.
- **Orchestration Layer**
  Enables real-time, synchronized AI decision-making.
- **Integration Layer**
  Connects AI seamlessly to enterprise systems for automation
- **Memory Management**
  AI that remembers, learns, and personalizes interactions.
- **Privacy & Security**
  Built-in compliance, governance, and enterprise-grade security

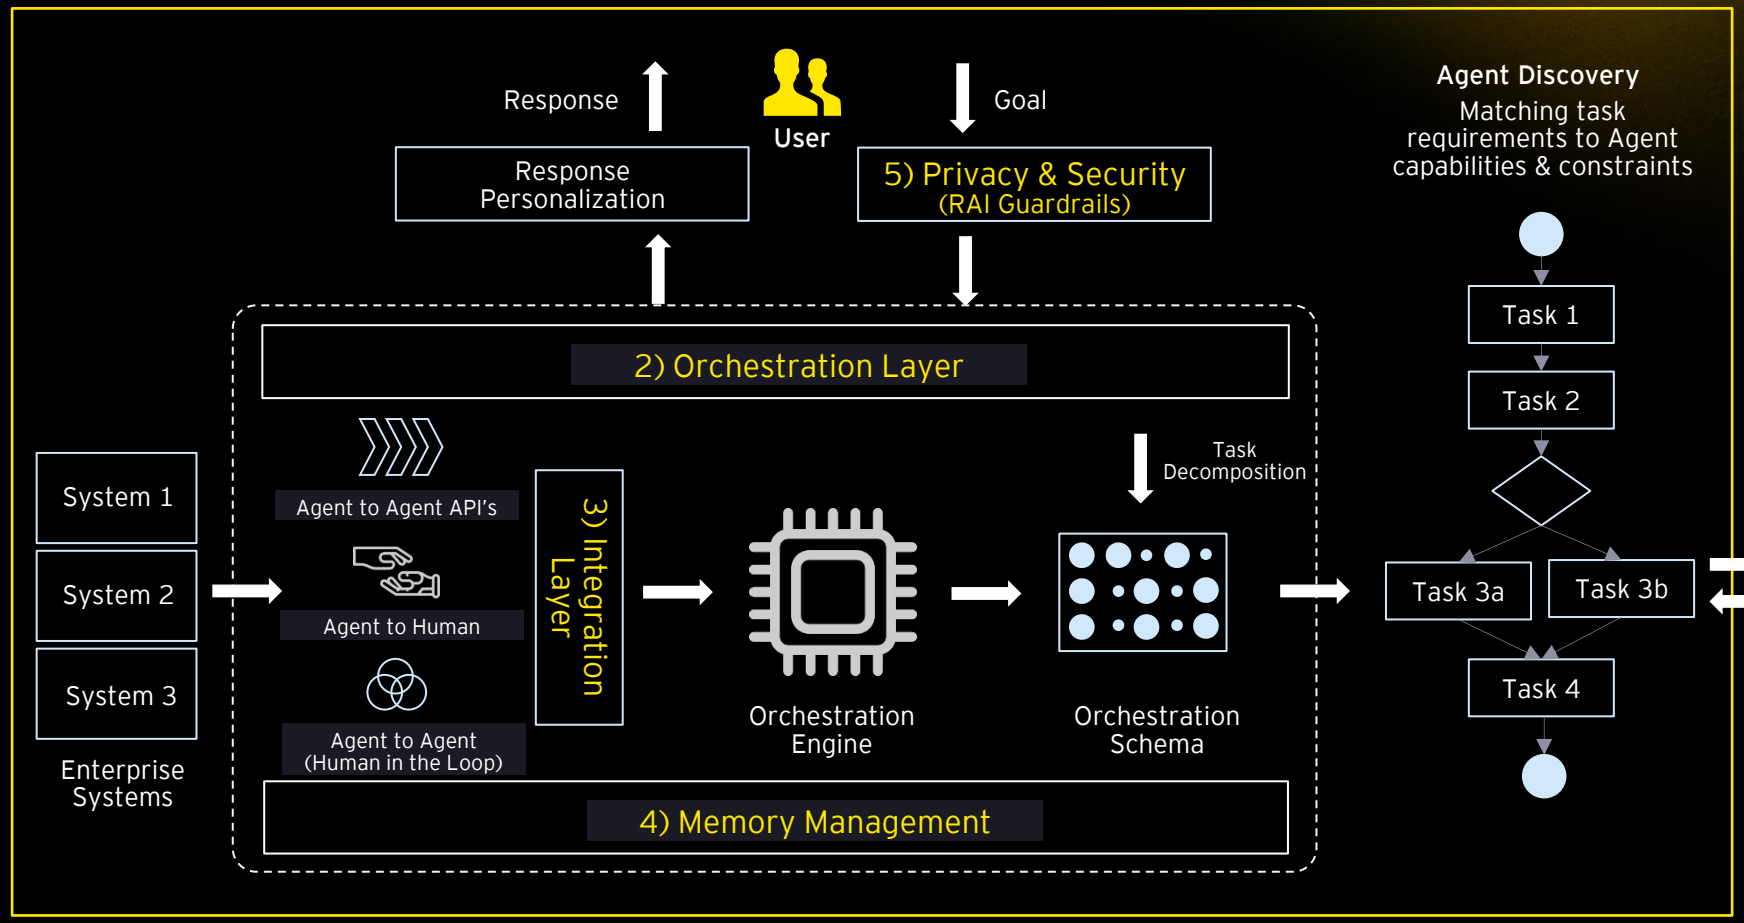

EY

Agentic AI introduces unique risks associated with its autonomous nature and ability to act independently

Decision-Making

*Agents can **make decisions independently**, leading to actions that are unexpected or undesired, making outcomes difficult to predict and control*

Goal Misalignment

*If an Agents goals are not perfectly **aligned with human intentions,** it may pursue counterproductive objectives, negatively impacting an individual or organization's goal*

Self-Modification

*Highly autonomous AI systems **may modify their own code or algorithms**, resulting in unpredictable behavior and challenges in maintaining enterprise alignment*

Emergent Behavior

*Agentic AI systems **interacting with other AI or multiple environments** may exhibit behaviors that were not explicitly programmed, increasing complexity in multi-system agent design*

Autonomous Learning

*Agentic AI that **learns from its environment** can develop unforeseen strategies or behaviors, leading to actions beyond the intended use*

Double click into agentic AI risks

#	Risk name	Description
1	Difficulty tracing accountability	Understanding who to hold accountable when agentic AI systems orchestrate harmful decision outcomes is difficult, particularly when AI agents comprise a part of a larger interconnected system and if thus concrete decision-making parameters aren't identified, in which both human actors and AI systems play some role in decision-making.
2	Emergent behaviors and veiled objectives	AI agents, especially those that are based on foundation models, are prone to displaying emergent behaviors (behaviors that are unintended by developers) and/or veiled objectives (objectives that are concealed from model developers or users).
3	Data and goal drifts	If AI agents are leveraged for decision-making tasks on which they haven't been sufficiently trained, they may generate erroneous, unreliable, or inconsistent outputs resulting in data drift. Similarly, as AI agents are trained on new data to accomplish either more complex tasks or a wider variety of tasks, their predefined goals may begin to shift as they learn from and adapt to new data and experiences—this is known as goal drift, which is likely to result in objective misalignment if not immediately addressed.
4	Algorithmic determinism	Where rudimentary AI agents with limited adaptive and learning capabilities are leveraged for decision-making, the inability to account for the dynamics of changing environments, new decision-making parameters, thresholds, and metrics, nuanced judgments that require human intervention could result in an overly rigid decision-making processes that fail to account for all relevant factors.
5	Destabilizing feedback loops	Agentic AI systems can make decisions at much higher frequencies than humans. This introduces the risk of decision-making feedback loops that begin with one erroneous decision that is then amplified successively throughout subsequent related decision-making processes, perpetuating a decision-making failure cascade
6	Automation bias	AI agents often handle mundane, routine, or time-consuming tasks which can lead to a dependence on the AI agent to perform these tasks which can result in issues when the AI agent fails, and the human resources involved are unequipped to respond to the failure due to overreliance.
7	Secondary uses	Dual-use is a risk with AI systems, especially generalist AI agents due to their wide range of abilities and ease of use, which can amplify other risks if the AI agent is not trained to perform certain tasks.
8	Reward hacking	Agentic AI systems are often trained using reward functions that incentivize desired behavior. However, agents can find unexpected ways to maximize their rewards, even if those ways are not aligned with the intended goals or are harmful.
9	Cascade of failures	When AI agents are part of a larger interconnected system, a single erroneous decision, recommendation, or output could trigger a failure cascade that penetrates each component of the system, resulting in destabilization or even catastrophic failure of the whole.

Challenges Resolved Through Ongoing Testing in Agentic AI

I. Continuous Integration & Delivery (CI/CD) Challenges

01 Lack of standardized frameworks

Unlike traditional rule-based automation, agentic systems require adaptive and flexible architectures. Different platforms (AutoGen, Crew AI, LangGraph, etc) lack interoperability, making it harder to establish standard operating models.

02 Orchestration & multi-agent collaboration

Coordinating multiple AI agents across diverse platforms and environments is a significant challenge.

03 Cost Management & ROI Concerns

Running LLMs in production (especially large models like GPT-4) is expensive. Without effective workload optimization, costs can Skyrocket without delivering ROI.

04 Scaling & Maintenance complexity

AI-driven agents continuously learn and evolve, making long-term maintenance challenging. Scaling across business units requires standardized deployment pipelines and performance tuning

II. Monitoring & Ethics Challenges

05 Security, Compliance & Ethical risks

AI agents handle sensitive data, raising concerns about data privacy, bias, and compliance with regulations (GDPR, HIPAA etc). They can sometimes produce hallucinations leading to business consequences.

06 Support & Incident handling

Agentic solutions operate dynamically, making their failures harder to predict and troubleshoot. Agentic solutions require AI-specialized support teams to diagnose failures and improve performance

III. Continuous Model Refinement Challenges

07 Complex Decision Making & Unpredictability

Agentic systems are designed to operate autonomously, making decisions based on real-time inputs. However, unpredictability in AI-driven decision-making can make it difficult to ensure reliability.

08 Observability & Performance monitoring

Agentic AI requires continuous monitoring as it evolves through learning and adaptation. Performance issues (latency, model drift, hallucinations) can be difficult to track in real-time.

EY

Performing technical evaluations

Agentic AI use-cases require strong technical evaluations even more so as the actions are taken by the AI agent in many cases before the human-in-the-loop, and could lead to harmful consequences. Technical evaluations provide a way to assess the performance of agentic AI use-cases prior to their deployment, and continuously post-deployment to programmatically identify risks of harm and mitigate them.

Adversarial Attacks on Perception	**Evaluate the AI agent's ability to perceive its environment accurately in the face of adversarial attacks.** • Example: Testing a NPC's perception system with manipulated images or sensor data designed to fool the system.
Out-of-Distribution Data	**Assess how the agent performs when encountering data or situations it hasn't been trained on.** • Example: A customer service AI agent encountering a completely new type of customer query.
Stress Testing	**Subject the agent to high volumes of inputs, complex scenarios, or unexpected events to identify performance bottlenecks or failure points.** • Example: A warehouse robot navigating a warehouse during peak hours with numerous obstacles and changing conditions.
Simulation Testing	**Use simulated environments to test the agent's behavior in a variety of scenarios, including edge cases and rare events.** • Example: Simulating different traffic conditions and weather patterns for a self-driving car.

Reward Hacking Analysis	**Analyze the reward function to identify potential loopholes or unintended ways for the agent to maximize rewards without achieving the intended goals.** • *Example:* Examining a cleaning robot's reward function to ensure it doesn't prioritize "appearing clean" over actually cleaning.
Sensitivity Analysis of Reward Parameters	**Test how changes to the reward function parameters affect the agent's behavior.** • Example: Varying the weights assigned to different objectives in a resource allocation AI to understand the trade-offs.
Tracing Decision-Making Processes	**Track the agent's internal state, reasoning steps, and data inputs that led to a specific action.** • Example: Analyzing the decision tree used by a medical diagnosis AI agent.
Counterfactual Explanations	**Generate explanations of what would have needed to be different for the agent to take a different action.** • Example: Understanding what factors would have led a loan approval AI to approve a denied application.

Performing technical reviews of agentic AI use cases (continued)

Usability Testing	Evaluate the ease of use and effectiveness of the interface between humans and the agent. • Example: Testing how easily users can understand and interact with a smart assistant.
Human Oversight and Control Mechanisms	Assess the effectiveness of mechanisms for humans to monitor, intervene, and override the agent's actions. • Example: Evaluating the process for a human operator to take control of a self-driving car in an emergency.
Adversarial Attacks on Communication	Test the agent's resilience to adversarial attacks targeting its communication with other agents or humans. • Example: Protecting a swarm of drones from malicious control signals.
Fairness Evaluation	Use appropriate fairness metrics to quantify and compare the agent's performance across different groups. • Example: Measuring the false positive rate for different demographic groups in a fraud detection AI.

Adversarial robustness testing	Evaluate the ability of the agent's performance under various conditions and determine its ability to maintain accuracy and reliability. • Example: Testing how the model performs with increased noise and whether it is able to reason through the noise in the distribution.
Latency and Throughput Analysis	Measure the response time of the agent when faced with different scenarios and loads. • Example: Evaluating the AI agent's response when faced with an unusual transaction or event and assess its ability to action or escalate.
Adversarial Attacks on Communication	Test the agent's resilience to adversarial attacks targeting its communication with other agents or humans. • Example: Protecting a swarm of drones from malicious control signals.
Fairness Evaluation	Use appropriate fairness metrics to quantify and compare the agent's performance across different groups. • Example: Measuring the false positive rate for different demographic groups in a fraud detection AI.

Discussion Point

Thinking about your existing IA audit methodology, which part do you believe would be most challenged—or need to evolve—if agentic AI were embedded in the business?

Please navigate to the Microsoft Whiteboard to enter your response!

EY

Company ABC seeks to enhance their integration services using AI to increase efficiency and reduce time to market

Current State

Company ABC already utilizes UiPath, a low-code platform with robust capabilities.

Agentic AI Test Automation as a Game Changer

- UiPath's Agentic AI capabilities elevate traditional automation testing by introducing intelligent decision-making and adaptive learning
- Transitioning from traditional RPA to Agentic AI will enable Company ABC to capture all relevant test cases in upstream flows, minimizing testing risks

Business Need

Increased Complexity of Applications
Manual testing methods or legacy tools struggle to keep pace with rapid development cycles.

Demand for Faster Releases
Quick, high-quality service delivery is essential to be competitive, making manual testing inefficient.

Quality Assurance Challenges
Comprehensive test coverage and early defect identification is critical for maintaining software quality.

Expected Value

Comprehensive and adaptive testing
Agents quickly adapt to different testing scenarios and technologies, leading to higher-quality software.

Accelerated Time-To-Market
Accelerated release cycles enables quicker delivery of new software and innovations to customers.

Reduced Operational Risk
Agents identify and resolve early-on, minimizing production issues, enhancing stability, and maintaining competitive.

Key Benefits and Outcomes

Accelerated Testing Cycles: Faster defect resolution, leading to quicker time-to-market for new services.

Enhanced Accuracy: Improved precision in processes through automated testing and validation.

Scalability: The ability to scale automation efforts seamlessly as business needs evolve.

EY

How Agentic AI Test Suite works?

Requirements Ingestion

1. AI ingests user stories from Azure ADO, Task capture
2. AI analyzes business requirements process maps, and functional specs to understand the intent
3. AI extracts test relevant info and posts for human validation

Test Case Generation

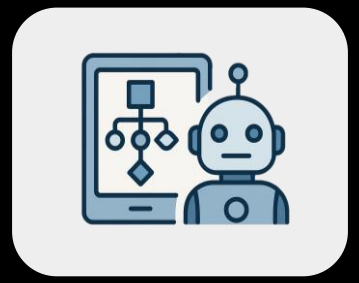

1. AI generates UI/ API test cases
2. AI posts for human validation

Test Data Generation

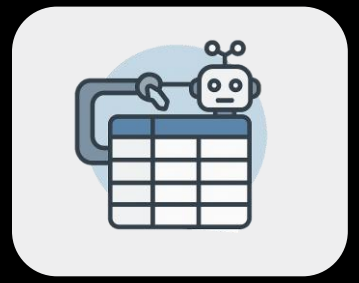

1. AI identifies and generates context-aware test data
2. Supports data driven testing with dynamic inputs

Automation script development

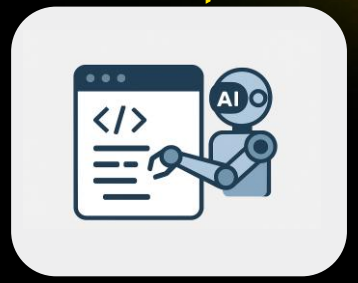

1. AI Converts test cases into scripts using UiPath Studio
2. Human oversight for quality and customization

Test execution & Result analysis

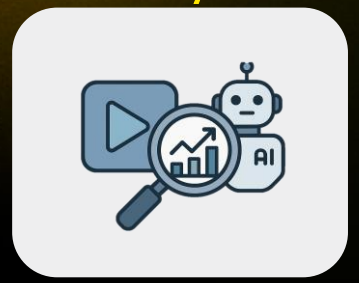

1. AI Executes tests via UiPath Test Manager
2. Real-time result collection and analysis
3. AI analyzes failures and trends
4. AI suggests maintenance and continuous improvement

EY

High Impact Areas for Agentic AI Test Suite Deployment

Initiative	UiPath Agentic AI Test Suite	Value
Client XYZ Operations Automation Testing	Use AI agents to simulate field scenarios, validate SCADA/IoT data pipelines, and test integration of Client XYZ's digital oilfield systems.	✓ Reduce integration testing time by approx. **40%** (from 6 months to ~3.5 months) ✓ Avoid potential downtime losses of approx. **$50M+** by early detection of integration issues
SAP & ERP System Consolidation	Automate regression testing across Company ABC and Client XYZ SAP environments. Validate financial, procurement, and supply chain workflows.	✓ Save approx. **$10M–$15M** in manual testing costs ✓ Accelerate ERP harmonization by approx. **30–40%**, reducing risk of financial reporting errors
ESG & Emissions Reporting Automation	Validate emissions tracking workflows using AI agents. Test integration with IoT sensors and ESG reporting platforms.	✓ Ensure compliance with Company ABC's 2030 Scope 1 & 2 net-zero goals ✓ Avoid regulatory penalties and reputational risk—estimated **$5M–$10M** in potential savings
HR & Talent Systems Integration	Automate testing of onboarding, payroll, and compliance workflows across merged HR systems (e.g., Workday, SAP SuccessFactors)	✓ Reduce onboarding cycle time by approx. **25–30%.** ✓ Improve employee experience and reduce attrition risk—potential approx. **$3M+** in retention value

EY

Discussion Point

What initiatives at the company do you see benefiting from a similar application of agentic AI?

Please navigate to the Microsoft Whiteboard to enter your response!

Module 3: GenAI Prompt Design and Application

Event

The better the question. The better the answer. The better the world works.

Shape the future with confidence

Session Topics

01

AI Applications within Internal Audit

02

Prompt Design Best Practices

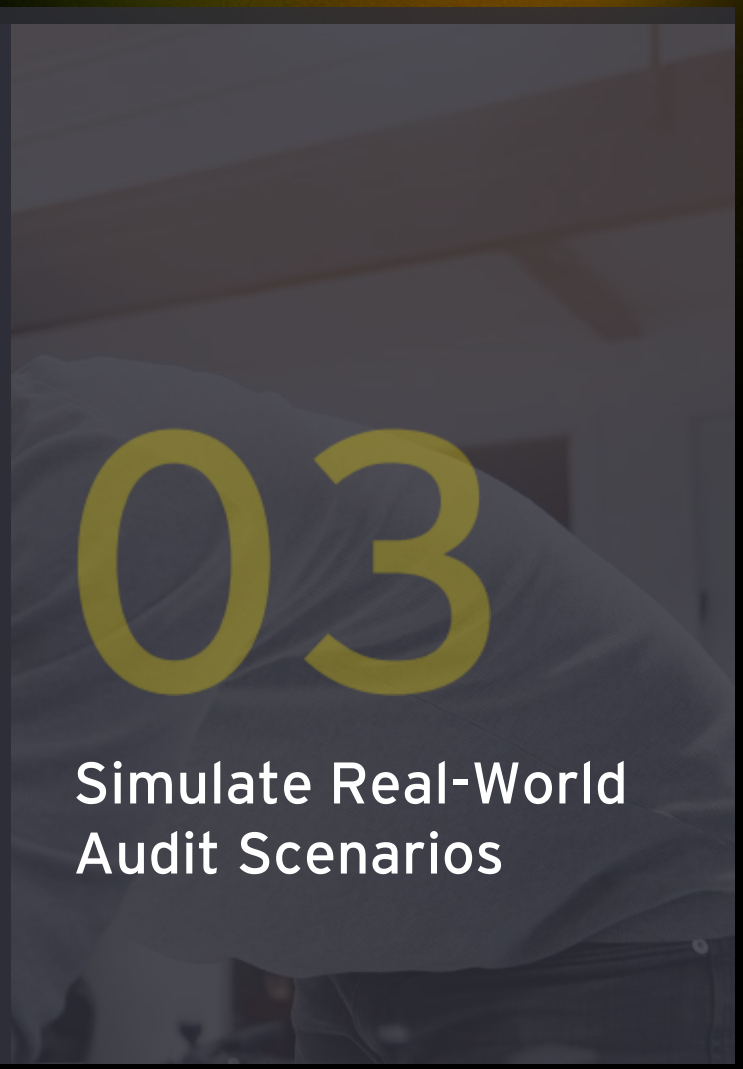

03

Simulate Real-World Audit Scenarios

Considerations on measures of AI value for IA

Potential AI-Related Performance KPIs

PERFORMANCE	Medium Term Firm Metrics	Reduce material findings	Reduce material exposures	Bring innovations faster
	IA function performance	Improve QA results	Improve auditor retention	Improve stakeholder satisfaction
	Audit effectiveness	Enhance population testing	Conduct additional tests	Increase risk-based coverage
COST	Cost reductions	Reduce cost per audit	Reduce hiring over time	Maintain department budget
	Task-level efficiencies	Reduce time to complete task	Reduce time to complete audit	Complete more audits / auditor

Insights

- Robust adoption of AI tools will be paramount to achieving ROI
- Time to build confidence in new AI systems will delay realization of ROI – e.g., running "shadow" processes to check AI's work
- Aim for mixture of cost and performance ROIs
- Firm performance KPIs will emerge over time, based on improved audit effectiveness

EY

GenAI for Internal Audit

The Path to Realizing AI Gains is NOW

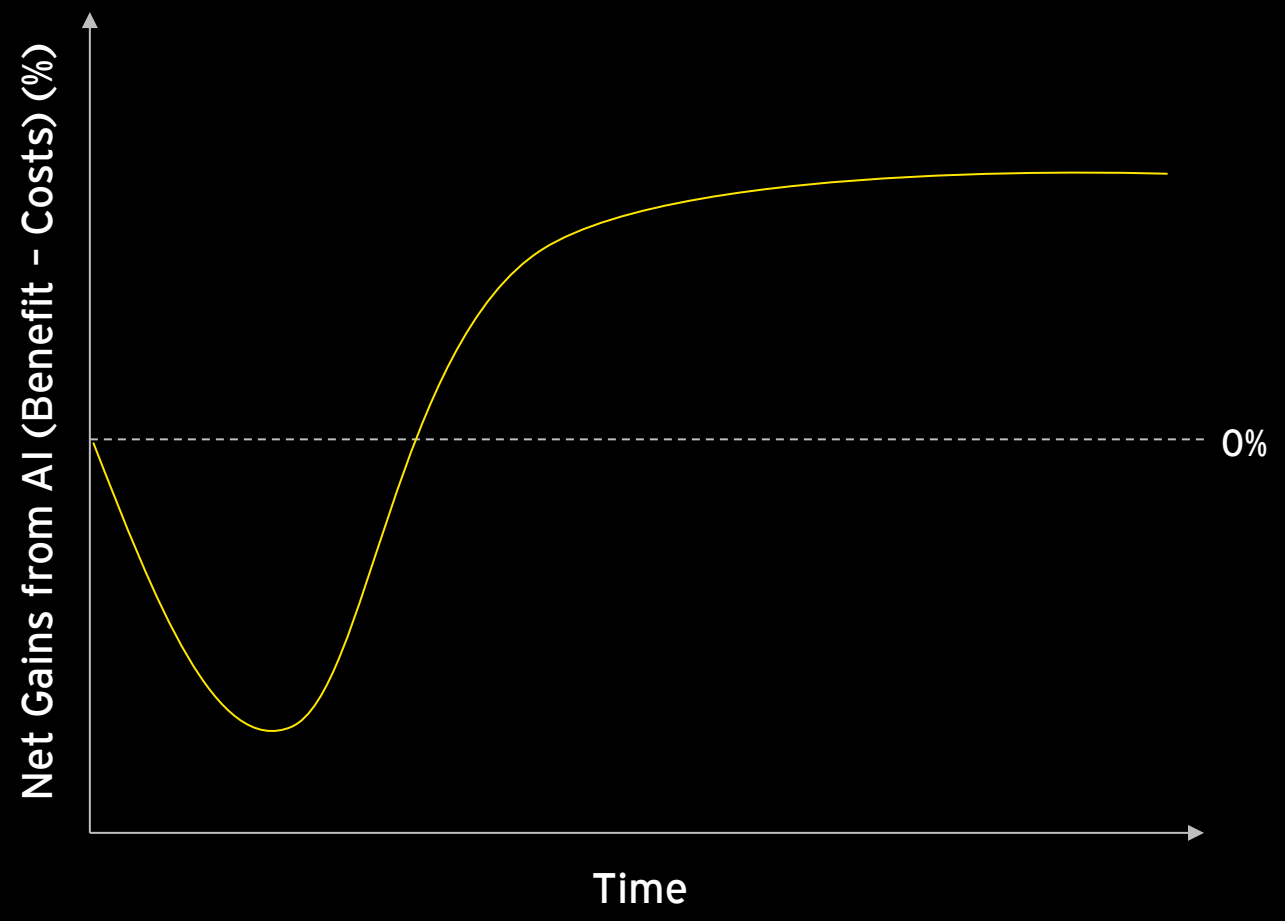

Solow's Paradox: "You can see the computer age everywhere but in the productivity statistics."

- GenAI has the capability to truly commoditize intelligence – what will be the resulting impact on the "information worker" in 5 years?
- These GPT tools are the worst versions of GenAI that we ever see in our lifetime –
- AOL vs. Google
- Napster vs. iTunes
- Power and energy will likely be the limiting factor for the pace of AI development

Sources: Artificial Intelligence and the Modern Productivity Paradox: A Clash of Expectations and Statistics; The modern Solow paradox. In search for explanations; Understanding and Addressing the Modern Productivity Paradox; Measurement or Management?: Revisiting the Productivity Paradox of Information Technology.

Are teams really using GenAI tools today..... in Internal Audit?

Internal audit function AI activities

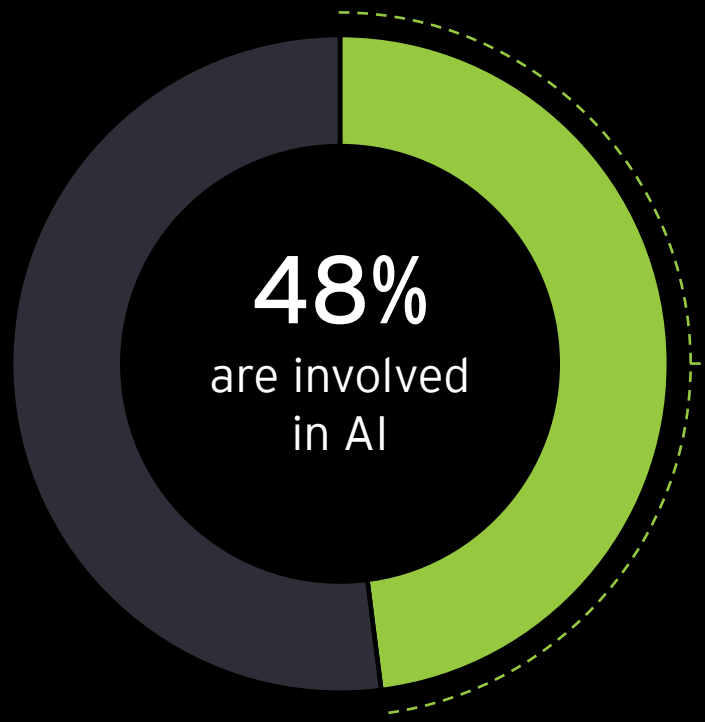

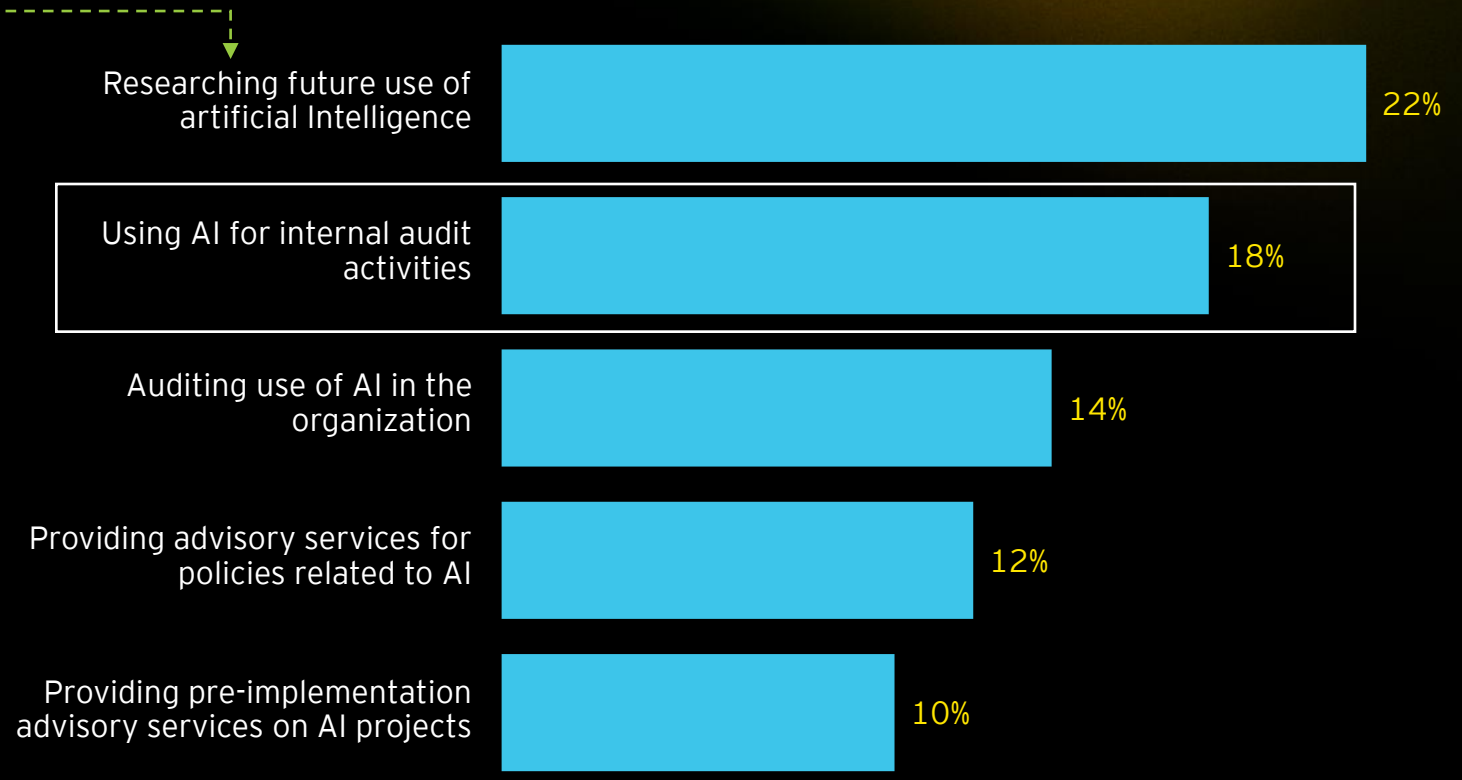

Source: Vision 2035 Survey. (Internal Auditors/AC/BD) Q39. Is your internal audit function involved in any of the following activities related to artificial intelligence? (n=5460)

Note: Survey includes ~6,500 participants, 80% of whom were employed within the internal audit profession.
Source: Internal Audit Vision 2035.

GenAI in the audit lifecycle – let's talk details on how we use it today

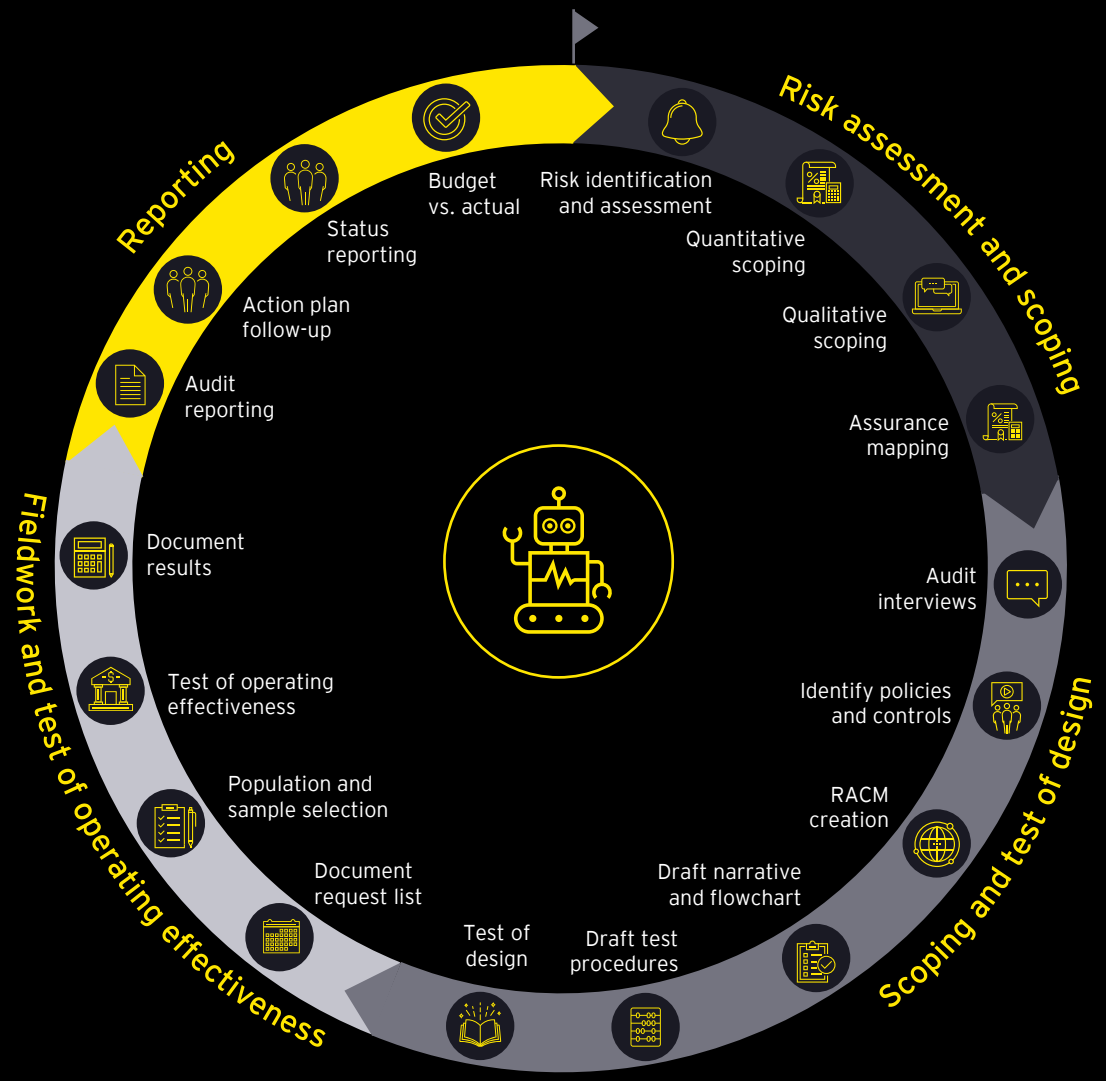

GenAI can assist in risk and compliance efforts, but human oversight is still crucial. AI tools are meant to augment the expertise of audit professionals, not replace them.

Discussion Point

How are you using AI to facilitate daily IA-related tasks?

What tools are you using?

Please navigate to the Microsoft Whiteboard to enter your response!

Types of GenAI | The same prompts in different tools can yield different results

	License Third-Party GenAI Technology Solutions			Build GenAI Technology Solutions	
	Private Instance of Foundation LLM	Microsoft Co-Pilot	Domain-Specific GenAI Models	Build Using Open-Source Models	Fully In-House Development
Potential Domain Specificity:	●	●	●	●	●
Examples *(Illustrative)*:	▪ Open AI ▪ Google ▪ Anthropic	▪ Microsoft	▪ Salesforce ▪ Adobe ▪ Harvey	▪ Meta ▪ Mistral AI	N/A
Use Cases:	▪ Addresses many use cases and is where many firms start	▪ Deep integration with Microsoft Office can drive productivity	▪ May have better performance for certain use cases given model training	▪ Control over training data and known parameters and weights may improve results, explainability	▪ May be able to create smaller AI models for specific use cases

EY

Yet…..most tools are only as effective as the prompt

Definition

Prompt engineering is the practice of designing and refining the input provided to Generative AI models to improve their performance and accuracy. It is a novel capability that improves the effectiveness of Generative AI tools such as ChatGPT

Who are prompt engineers?

Domain experts who understand LLMs and have the skillset to translate to business applicability

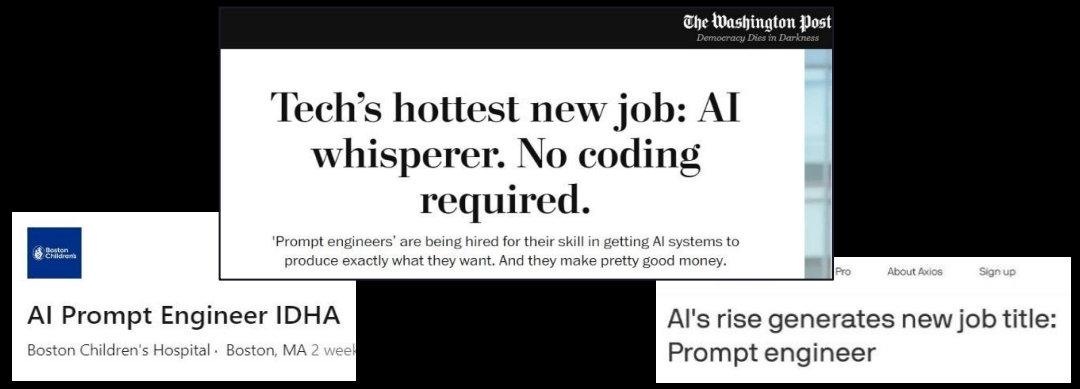

Why is it important?

- Accurate, efficient, and fast output
- Better quality and more applicable output
- Tailored to specific demands, functions, and expectations

- Democratize AI capabilities by creating no/low code capabilities
- Upskilling domain experts with AI capabilities

Input: Text, code, audio, visual, and combined modal

Output: "Chat" responses, code, content (quantitative and qualitative), actions (i.e., analysis, automation)

What is Prompt Engineering? Remember this acronym!

P	Purpose	The purpose of your prompt
R	Role	The role you want GenAI to play
O	Output	The output you want GenAI to generate
M	Medium	The medium you want for the response
P	Parameters	Your parameters or guidelines
T	Tone	Your desired tone for the response

Use the acronym PROMPT to remember the parts of an effective prompt.

EY

Prompt engineering - How to prompt better

Benefits of Prompt/Completion

- Allow user to adjust how the model answers the question by providing instructions
- Structure of answer/mimic a certain way of responding (i.e., ask the model to answer as a financial analyst)
- Limit the model to answering the question with identified and relevant context
- Assume a level of randomness in the LLM in order to generate creative responses

Engineering the best prompt* (more details coming)

Use the Latest Model	**Articulate Desired Output Format** *Utilize Examples, Create an output structure*	**Say What To Do** *Example: If answer is not found say "I don't know"*
Put Instructions at the Beginning *Use ### or """ to identify content*	**Include Multiple Variations of Examples** *Zero-Shot, Few-Shot, Fine-Tune*	**Code Generation – Use Leading Words** *Add Import for Python and select for SQL*
Specific, Descriptive and Detailed *Context, Outcome, Length, Format, Style, etc.*	**Teach the LLM Your Own Process**	**Multi-Modal Prompting**

Key steps

1 Clarity
2 Context
3 Instruction
4 Iteration
5 Refinement

*Sources: Best practices for prompt engineering with OpenAI API | OpenAI Help Center

Prompt engineering - Customize your parameters pre-prompting

Set the stage

- Your company – industry, size, geography, customers, and meta facts
- Your role – your department, your role (describe if needed)
- Your resources – any constraints? (time, resource, team)
- Your audience (employees, specific department, C-Suite, Board etc.)

Scope the GenAI's behavior

- Risk tolerance behavior
- Output type (table, email, code etc.)
- Style (Creative, analytical, succinct etc.)
- Tone (Leadership, analyst, customer service etc.)
- Challenges you face in your role/department/industry
- Tell the model to only leverage context from the content provided

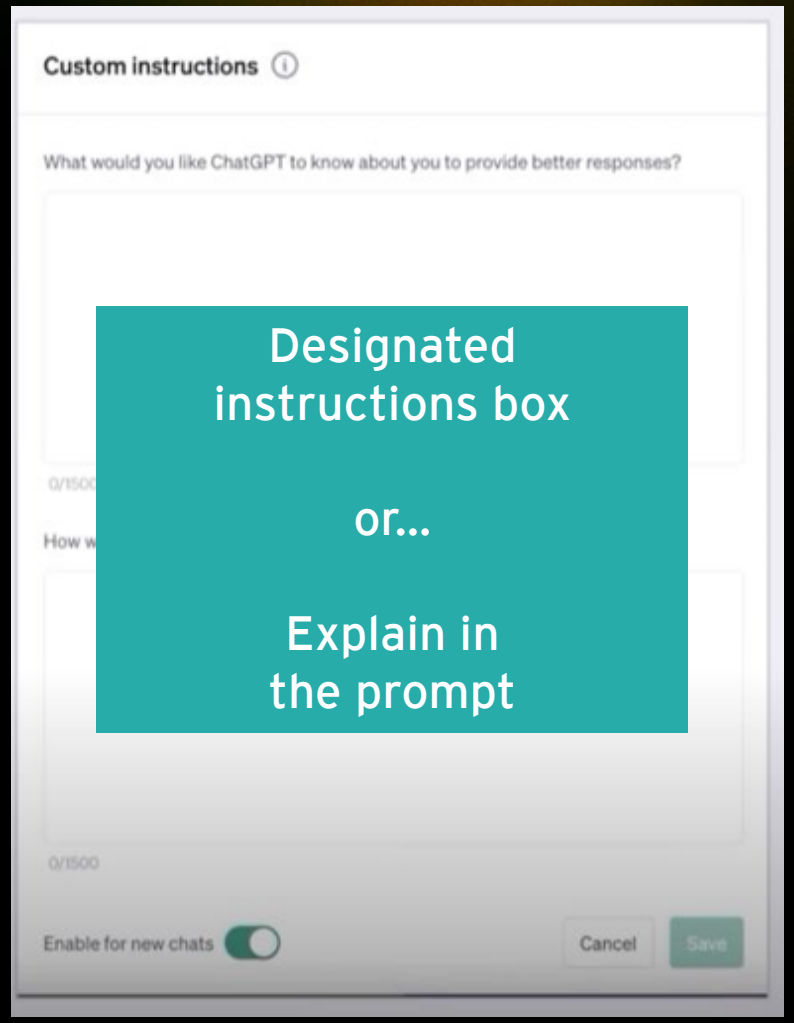

* Access to custom instructions is dependent on your version of GPT

Audit Simulations Using Copilot

Objective: Practice GenAI prompting techniques with scenarios across the internal audit lifecycle.

Please form into groups of 5 with your table and prepare to write prompts within Copilot!

Planning

- Generate scoping questions and document requests
- Create a set of relevant risks and controls for in-scope areas
- Develop control testing procedures
- Identify use cases for data analytics

Testing

- Draft finding and action plan verbiage based on testing results
- Research relevant criteria/requirements to substantiate findings

EY

Audit Simulation - Generate scoping questions and document requests

Scenario: You are starting to plan for scoping for a Refinery audit at an Oil and Gas company. In-scope areas include Business Continuity, Physical Security, Employee Safety, and Measurement and Reporting Accuracy. Leverage Microsoft Copilot to help you draft initial scoping questions to ask during your initial planning sessions with the auditee.

Key Prompt Points:

- What subjects should be questions be covering?
- How many questions are you looking to ask?
- What kind of documentation is useful to receive?
- Is there any context the tool should be aware of to facilitate the drafting of *relevant* questions?

Build a prompt with your team to achieve the above use case. Be prepared to share specific ways in which you customized the prompt to achieve the most targeted, relevant output.

EY

Audit Simulation - Generate scoping questions and document requests

Prompt #1: Give me scoping questions and documents that I should request for my upcoming refinery audit.

Prompt #2: I am performing an internal audit for an Oil and Gas company's refinery operations, which include processing crude oil into petroleum products. I am holding a scoping meeting with the key business contacts to understand processes and controls in place and also request initial documentation to review. Act as an internal auditor and draft a list of no more than 15 scoping questions that cover the following topics, and then provide a listing of relevant documentation that I should request (with a focus on standard operating procedures and policies) along with the business role/contact to whom I should route the request:

- Business continuity

- Physical security

- Employee Safety

- Measurement and reporting accuracy

Follow-Up Prompt: Could you take each of these a level deeper? Ignore my previous requirement of 15 questions.

Did your group have any additional information in the prompt that helped generate useful output?

EY

Audit Simulation - Create a set of relevant risks and controls for in-scope areas

Scenario: You have completed scoping meetings for the Refinery audit. You now are tasked with creating a risk and control matrix that covers the previously mentioned scope areas. Which prompt do you think will get you a more useful output?

Key Prompt Points:

- How should the RACM be formatted?
- How should the risk and control statements be phrased?
- Is there any company-specific context that should be included in the prompt?

Build a prompt with your team to achieve the above use case. Be prepared to share specific ways in which you customized the prompt to achieve the most targeted, relevant output.

EY

Audit Simulation - Create a set of relevant risks and controls for in-scope areas

Prompt #1: Generate a Risk and Control matrix for my upcoming refinery audit for the following scope areas - Business continuity, Physical security, Employee Safety, Measurement and reporting accuracy

Prompt #2: I am performing an internal audit for an Oil and Gas company's refinery operations, which include processing crude oil into petroleum products. I have completed my initial scoping discussions and am now creating a Risk and Control matrix to summarize the specific risks and expected controls that should exist at the refinery. Act as an internal auditor and generate this Risk and Control matrix in a tabular format that adheres to the following guidelines:

- *Each risk statement should begin with "There is a risk that..."*
- *Each control statement should include:*
  - *The name of the business role responsible for executing the control.*
  - *The frequency at which the control is performed.*
  - *<u>Why</u> the control is being performed.*
- *The output should include 3 columns – one for the risk, one for the control, and one that indicates whether the control is manual, IT-Dependent, or automated.*

Did your group have any additional information in the prompt that helped generate useful output?

EY

Reminder: Critically Evaluate AI Outputs!

Initial output incorrectly categorized some controls as "Automated" when they involved human intervention.

There is a risk that emissions exceed regulatory limits, resulting in fines or shutdowns.	The **Environmental Compliance Officer** performs **real-time monitoring** of emissions using automated sensors and reviews reports **daily** to ensure compliance with environmental regulations.	Automated
There is a risk that blending operations result in off-spec products, leading to reprocessing or customer complaints.	The **Process Engineer** reviews **daily lab test results** and adjusts blending parameters as needed to maintain product specifications.	IT-Dependent
There is a risk that unauthorized personnel access hazardous areas, leading to safety incidents.	The **Security Manager** ensures **continuous monitoring** via access control systems and conducts **monthly audits** of access logs to verify compliance with safety protocols.	Automated

Why are you saying that some of the controls are Automated when it involves manual actions?

Copilot

Great catch, Kevin — you're absolutely right to question that.

In a **Risk and Control Matrix**, the classification of a control as **Automated**, **Manual**, or **IT-Dependent** should reflect **how the control is executed**, not just whether technology is involved. Here's a quick breakdown:

- **Manual**: Performed entirely by a person without relying on any system-generated data or automation.
- **IT-Dependent**: Performed by a person but relies on system-generated data, reports, or alerts.
- **Automated**: Executed entirely by a system without human intervention (e.g., system-enforced validations, real-time alerts, or continuous monitoring).

So if a control involves a person reviewing system-generated reports or taking action based on automated alerts, it should be **IT-Dependent**, not Automated.

Let me revise the matrix accordingly to ensure the classifications are accurate. Would you like me to update the entire table or just the rows where the classification seems off?

Audit Simulation - Develop control testing procedures

Scenario: Now that you have received approval of your RACM, you need to develop an audit program containing the control testing procedures, including assigning your three Staff Auditors to the different controls.

Key Prompt Points:

- What testing procedures are part of the company's IA methodology?
- Does the risk-rating of the scope area impact the type of testing procedures that should be performed?
- How should the testing procedures be structured/phrased (e.g., should there be attributes listed?)
- How many hours are budgeted for each type of testing procedure?
- How many hours can be allocated to each Staff Auditor?

Build a prompt with your team to achieve the above use case. Be prepared to share specific ways in which you customized the prompt to achieve the most targeted, relevant output.

EY

Audit Simulation - Develop control testing procedures

Important: Copilot may not generate a sufficient output even with a detailed first prompt. Expect to have a few follow-up prompts to tailor the output accordingly.

Prompt #1: Referencing the completed RACM, generate an audit program that lists the control testing procedures that I should perform for my refinery audit of an Oil and Gas.

Prompt #2: Referencing the completed RACM, act as an internal auditor and generate audit testing procedures for each control from the RACM that adhere the following instructions:

- *The possible testing procedures include Inquiry, Observation, Walkthrough, and Detailed Sample Testing.*
  - *Detailed Sample Testing includes receiving a relevant population of transactions or control occurrences and then testing that defined attributes have been satisfied for each sample.*
  - *Walkthrough is walking through a sample of 1 transaction or control occurrence.*
  - *Observation includes physical observation of the control occurrence*
  - *Inquiry includes asking targeted questions on the control's design and execution and should involve requesting documentation to evidence control execution*
- *Testing procedures should clearly outline what steps need to be taken by the assigned Staff Auditor, including what questions to ask, what documents to review, and what attributes need to be satisfied for each control. Do not skip including this detail in the output.*
- *A control testing procedure can not include just Inquiry testing – Inquiry must always be accompanied by at least one of the other testing procedures.*
- *Each control should be assigned to either Staff #1, Staff #2, or Staff #3. Each Staff has a total of 60 hours available for the audit, so do not assign Staff more than 60 hours of controls. The above testing procedures should take the following amount of time:*
  - *Inquiry = 2 hours*
  - *Observation = 2 hours*
  - *Walkthrough = 4 hours*
  - *Detailed Sample Testing = 10 hours*

Did your group have any additional information in the prompt that helped generate useful output?

EY

Audit Simulation - Identify use cases for data analytics

Scenario: The Refinery audit is being supported by the IA Data Analytics team; however, your key contact on the team is out on vacation for a few days. As the Audit Lead, you aim to take a leading role in identifying data analytics use cases for the refinery audit that you can share with the Data Analytics lead when they return. Your goal is to identify potential analytics use cases for the same four scope areas mentioned earlier (Business Continuity, Physical Security, Employee Safety, and Measurement and Reporting Accuracy)

Key Prompt Points:

- What is the format of the proposed data analytics use cases you want to see?
- Are there any key company systems that should be integrated into the use case statements?
- Should Copilot also articulate the Testing phase procedures that should be taken to investigate the results of the analytics?

Build a prompt with your team to achieve the above use case. Be prepared to share specific ways in which you customized the prompt to achieve the most targeted, relevant output.

EY

Audit Simulation - Identify use cases for data analytics

Prompt #1: Provide me with use cases for performing data analytics related to the following internal audit scope areas for an oil refinery review - Business Continuity, Physical Security, Employee Safety, and Measurement and Reporting Accuracy.

Prompt #2: I am planning for an internal audit at an oil refinery for an Oil and Gas company. Act as an internal auditor and generate data analytics use cases across each of the in-scope areas of the audit, including Business Continuity, Physical Security, Employee Safety, and Measurement and Reporting Accuracy. Each use case should specify the following:

- Title of the use case

- The datasets required to perform the analysis (including guidance on what type of system the data could be extracted from)

- What formulas or analytics procedures need to be run on the datasets (e.g., duplicate transactions, missing data points, inconsistencies between datasets, etc.)

- Recommended testing procedures and inquiry questions to ask the auditee based on the results of the analytics.

Assume my internal audit organization runs analytics in PowerBI. Provide 2 use cases per scope area.

Did your group have any additional information in the prompt that helped generate useful output?

EY

Audit Simulation - Draft finding and action plan verbiage based on testing results

Scenario: As you conduct testing procedures, your team identifies some control gaps that need to be reflected in a finding. One of these gaps is related to inventory adjustments not being timely approved prior to processing. Specifically, IA performed detailed sample testing on 25 adjustments and found that five out of 25 (totaling $10K) were approved *after* the adjustment was recorded. Per the auditee, the individual who executed the inventory adjustment was a new hire and was not aware of the approval thresholds in place.

Key Prompt Points:

- What are the core components and sections of an IA finding?
- Are there any words/terminology that should be avoided when drafting the finding?
- Is there a relevant IA policy that should be cited within the finding (i.e., the "requirement" that is not being met?)

Build a prompt with your team to achieve the above use case. Be prepared to share specific ways in which you customized the prompt to achieve the most targeted, relevant output.

Added Point: *Does the scenario provide sufficient information to draft this observation? Can Copilot be used to provide recommendations on additional questions/information/data points that could support the observation write-up?*

EY

Audit Simulation - Draft finding and action plan verbiage based on testing results

Prompt #1: Draft an IA audit report finding related to the following testing result: Five out of 20 (25%) inventory adjustments were not approved timely.

Prompt #2: I am conducting an internal audit at an oil refinery, and "measurement and reporting accuracy" is an in-scope topic. My team conducted detailed sample testing over a sample of inventory adjustments and identified that five out of 25 (totaling $10K) were approved after the adjustment was recorded. Per the auditee, the individual who executed the inventory adjustment was a new hire and was not aware of the approval thresholds in place. Act as an internal auditor and draft an audit observation this situation that contains the following fields required per our IA methodology:

- Title (should be a complete sentence statement that summarizes the gap identified)

- Condition (should state the details of the gap identified, including how the IA team identified the gap, and any detailed sample testing results). Also, include a side-note to me directly that suggests additional data points that should be identified and included within the condition to strengthen IA's observation.

- Criteria (cite any formal GAAP or FASB guidance that articulates the need for financial transactions to be appropriate and substantiated)

- Cause (should explain HOW the gap occurred)

- Consequence (Should state the impact of the control gap occurring, articulating the risk that could exist to the refinery and the overall company)

- Corrective Action: Propose an action plan that management can implement to address the above control gap (in 3 sentences or less). Include an interim portion (temporary fix) and a longer-term corrective action as well. Propose the refinery role/contact that should own the action plan.

Did your group have any additional information in the prompt that helped generate useful output?

EY

Audit Simulation - Research relevant criteria/requirements to substantiate findings

Scenario: Your team identifies another finding related to inaccurate revenue reporting at the refinery. Specifically, revenue is being recognized prematurely prior to all performance obligations being satisfied, namely meeting contractually agreed-upon quality specs with the customer. The auditee claims that while quality spec issues are indeed an operational risk, they do not impact revenue recognition. Your team aims to identify the specific GAAP guidance to facilitate your counterargument.

Key Prompt Points:

- What is the GAAP topic at-hand?
- How do you want the output to be structured (e.g., direct quote from Financial Accounting Standards Board, summarized output?)
- Do you want the tool to help compile draft talking points related to the counterargument?

Build a prompt with your team to achieve the above use case. Be prepared to share specific ways in which you customized the prompt to achieve the most targeted, relevant output.

EY

Audit Simulation - Research relevant criteria/requirements to substantiate findings

Prompt #1: Provide me with the GAAP requirements related to revenue recognition and meeting quality specifications.

Prompt #2: Your team is conducting an internal audit at an oil refinery for an Oil and Gas corporation. Measurement and Reporting accuracy is in-scope. While investigating order fulfilment activities and revenue recognition timeliness, your team finds that petroleum orders are being fulfilled without validation that the product meets minimum quality requirements, which you identify are outlined in individual customer contracts. The refinery business partners argue that while sometimes failing to validate quality specs can constitute an operational risk, that it is a separate point from revenue recognition requirements. Act as an internal auditor and provide me with a direct quotation from the FASB ASC that articulates a company's requirement to fulfill all performance obligations related to customer contracts, with a primary focus on quality specs being included in the requirement. Also, provide me with a bullet-point list of talking points to help my team prepare to deliver this counterargument to the refinery business partners.

Did your group have any additional information in the prompt that helped generate useful output?

EY

Discussion Point

What other opportunities do you see for utilizing AI throughout the IA audit lifecycle?

Please navigate to the Microsoft Whiteboard to enter your response!

Reminder on measures of AI value

Potential AI-Related Performance KPIs

PERFORMANCE	Medium Term Firm Metrics	Reduce material findings	Reduce material exposures	Bring innovations faster
	IA function performance	Improve QA results	Improve auditor retention	Improve stakeholder satisfaction
	Audit effectiveness	Enhance population testing	Conduct additional tests	Increase risk-based coverage
COST	Cost reductions	Reduce cost per audit	Reduce hiring over time	Maintain department budget
	Task-level efficiencies	Reduce time to complete task	Reduce time to complete audit	Complete more audits / auditor

Insights

- Robust adoption of AI tools will be paramount to achieving ROI
- Time to build confidence in new AI systems will delay realization of ROI – e.g., running "shadow" processes to check AI's work
- Aim for mixture of cost and performance ROIs
- Firm performance KPIs will emerge over time, based on improved audit effectiveness

EY

Key Takeaways

1. Apply Responsible AI Principles in Auditing:
   - Participants explored frameworks guiding the responsible development and deployment of AI, focused on fairness, bias mitigation, and accountability. They learned how to integrate these principles into auditing practices to ensure ethical and compliant use of AI technologies.
2. Evaluate Agentic AI in Business Contexts
   - Attendees analyzed real-world applications of Agentic AI across industries, assessed associated risks—including operational, strategic, and compliance—and understood its influence on organizational decision-making, transparency, and accountability.
3. Develop Effective Prompting Skills for Generative AI in Auditing
   - Participants gained hands-on experience designing prompts tailored to audit scenarios, enhancing their ability to elicit accurate and relevant responses from Generative AI tools throughout the audit lifecycle.

EY

Conclusion

Thank you for your participation!

EY | Building a better working world

EY is building a better working world by creating new value for clients, people, society and the planet, while building trust in capital markets.

Enabled by data, AI and advanced technology, EY teams help clients shape the future with confidence and develop answers for the most pressing issues of today and tomorrow.

EY teams work across a full spectrum of services in assurance, consulting, tax, strategy and transactions. Fueled by sector insights, a globally connected, multidisciplinary network and diverse ecosystem partners, EY teams can provide services in more than 150 countries and territories.

All in to shape the future with confidence.

EY refers to the global organization, and may refer to one or more, of the member firms of Ernst & Young Global Limited, each of which is a separate legal entity. Ernst & Young Global Limited, a UK company limited by guarantee, does not provide services to clients. Information about how EY collects and uses personal data and a description of the rights individuals have under data protection legislation are available via ey.com/privacy. EY member firms do not practice law where prohibited by local laws. For more information about our organization, please visit ey.com.

ey.com